

MO444 Final Assignment - Beer Style Classification

Pedro Henrique M. X. Zacarin*

Abstract

In this project, a dataset with nearly 75000 homemade beer recipes was utilized in order to train a beer style classifier. The dataset contains data of 170 different beer styles with a total of 5 useful features and is heavily imbalanced in its examples-per-class quantities. Allied to the fact that many beer styles shares common attributes, classification may be difficult and may not provide great accuracy results.

1. Introduction

The production of homemade beers was always something important between the drink lovers. Many prefer making its own beer due to a set of reasons, such as spending less money with industrialized ones, picking specific ingredients, appeal to certain groups of people with different tastes or just for experimentation, to make something new.

The Internet has helped homebrewers's community a lot, as it enabled recipes and manufacturing techniques sharing and the opportunity to buy and sell different, and many times exotics and imported, ingredients. Furthermore, it allowed beer lovers to learn more about the production process, first steps, materials, ingredients and more intricate subjects involving chemistry and ingredient interaction.

Each beer has its style, which is a term used to differentiate and categorize them by various factors such as color, appearance, production method, alcoholic content, origin or even history. There is a long list of styles, ranging from classic ones such as Bock, Altbier and Pale Ale to specialty styles such as Fruit Beer and Honey Beer.

In the website Kaggle (www.kaggle.com, a popular data science website that offers a great number of datasets) a dataset of nearly 75000 homebrewed beers of over 170 different styles originally shared in the website Brewer's Friend (www.brewersfriend.com) was available in order to be utilized for analysis, classification or sheer curiosity. That dataset was utilized in this project in order to try to assess the correlation between the beers features and its styles by building a classifier that predicts the beer style based on its features.

*Contact: phzacarin@gmail.com

2. Dataset

The original dataset needed to have its information treated in order to be utilized. First, the features which didn't have important information for classification or that had a great amount of missing data were removed, such as priming method, pitch rate and priming amount. Then, all rows corresponding to an original gravity and final gravity measured in Plato (measurement of the concentration of dissolved solids in a brewery wort) were also removed, as the vastly majority of the examples in the dataset were measured in another scale: Specific Gravity. The final treatment was to eliminate all beer styles with just one example, as it wouldn't be possible to separate a training and test set with each having at least one example of every class.

After the treatment mentioned above, the final dataset contained 67768 examples from 168 distinct styles with 5 features each that could be utilized in order to train the classifier: original gravity (OG), final gravity (FG), alcoholic content (ABV), bitterness (IBU) and color which are all important measures of the brewing process. The color feature was already encoded into a continuous scale known as SRM (Standard Research Method).

3. Proposed solutions

As mentioned in prior sections, the dataset utilized is heavily imbalanced. To illustrate such characteristic, a histogram with the frequency of each style was generated (Figure 1).

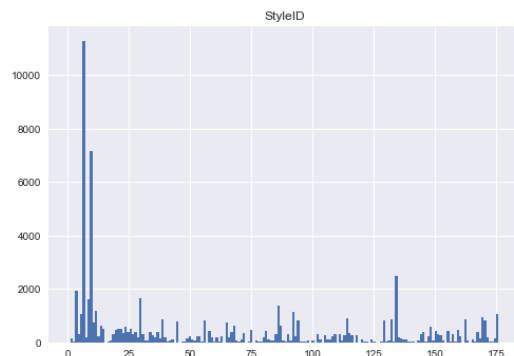


Figure 1. Histogram showing style frequencies

The three most predominant styles have frequencies of 11270 (American IPA), 7170 (American Pale Ale) and 2466 (Saison) examples each. It's worth noting that the 11270 examples of American IPA style corresponds to roughly 1/6 of all examples and the top 13 most predominant styles corresponds together to half of all examples.

At first, the imbalance will be disregarded as an important characteristic and the whole dataset will be considered in order to attempt to train a classifier. A set of classification algorithms are going to be used to assess each one's resulting accuracy, with 80% of the dataset for training and 20% for cross validation. The algorithms are the following:

- K-Nearest-Neighbors (K-NN)
- Support Vector Machines (SVMs) with One vs. One approach
- Random Forests (RF)

After each one of these methods were applied, a better approach will be considered: splitting the dataset into two, training the first one with the majority of the examples with one classifier and the other one, with classes with fewer examples, with another one. For that second scenario, a separation of 80%/20% for training and cross validation sets will be utilized to evaluate performance. K-fold validation in this case can't be applied because of the small number of examples in some classes, which makes a stratified division in K folds of the dataset impossible.

4. Development

The first attempt made in order to create a classifier was to consider the whole dataset for the classification. The training set and test set were initially created considering the number of examples of each style, which leads to both sets preserving the percentage of samples for each class. For that scenario, 80% of the dataset was utilized for training and 20% for validation purposes.

K-Nearest-Neighbors is an algorithm utilized for both classification and regression based on feature similarity, where an object is classified by a majority vote of its neighbors, with the object being assigned to the class most common among its K nearest neighbors. To choose parameter K, grid search was performed (Table 1).

K-Nearest-Neighbors algorithm resulted in training and cross validation accuracies of 34.86% and 32.59% respectively for K = 50, which is far from a great result.

The next method applied was Support Vector Machines, which tries to find the best separation between classes by finding an hyperplane between the support vectors of each class considered. As the project's problem is multiclass,

K	Training accuracy (%)	Validation accuracy (%)
10	39.93	30.46
20	37.19	32.04
30	36.11	32.43
40	35.37	32.56
50	34.86	32.59
60	34.46	32.25
70	34.14	32.36
80	33.89	32.23
90	33.58	32.00
100	33.48	31.92
150	32.75	31.86

Table 1. Accuracies for each value of K in K-NN algorithm for whole dataset classification

one vs. one approach was utilized, which means that comparisons between classes will be made two-by-two for every pair available. The training and validation accuracies obtained were respectively 39.12% and 33.42%.

As a last try with the first scenario, Random Forest was utilized, which is an ensemble learning method for classification and regression that operate by constructing many decision trees at training time and utilizing bagging for generating them. Grid search was performed to assess the best number of estimators (number of trees in the forest) to be used (Table 2). The parameters for Random Forests algorithm were thoroughly tested in order to find the best combination avoiding overfitting, which are the following:

- Number of features to be considered when looking for the best split: square root of the number of features
- Minimum number of samples required to split an internal node : 4
- Maximum depth of the tree: 15

In order to find a good number of estimators (trees in the forest), grid search was performed and the following results were obtained:

N. estimators	Training accuracy (%)	Validation accuracy (%)
50	57.88	35.92
100	58.34	36.03
200	58.44	35.92
300	58.61	35.91
400	58.60	35.93
500	58.59	35.93

Table 2. Accuracies for each number of estimators (trees) in Random Forests algorithm for whole dataset classification

As grid search shown, there is not a substantial difference between the number of estimators and validation accu-

racy, so 200 was considered. With it, training and validation accuracies obtained were respectively 58.34% and 36.03%.

In order to try to improve both training and validation accuracies, a split in the dataset was performed and each part of it was trained with a different method in order to mitigate, at least slightly, the effects of heavy class imbalance.

To find the best split for the original dataset (examples frequency threshold), grid search was utilized. For the first part of splitted dataset, which contains classes with more examples, Random Forests algorithm was used with 200 estimators, as it was the one with better results in the previous scenario and because it works great with more examples. For the second part, which contains classes with fewer examples, SVMs with one vs. one approach and Random Forests will be employed to assess which one performs better in a case with fewer examples per class. A 80% and 20% separation between training and validation for each one of the parts from the original dataset was made. The results are displayed in Table 3 and Table 4.

Split frequency	Examples (1st half)	Classes (1st half)
200	61767	79
300	56949	59
400	50943	41
500	46504	31
550	44967	28
600	43791	26
650	41909	23
750	40420	21
800	39631	20
850	35464	15
900	33729	13
950	31885	11

Table 3. Number of examples and number of classes in first part of the splitted dataset for each frequency point of split

The results shown something that was already expected: the fewer classes in the first half dataset, the better the training and validation accuracy utilizing Random Forests. In order to find a sweet spot, it is necessary to run SVMs and Random Forests algorithms for the second half for the same frequency splits. The results can be seen in Table 6 and Table 7.

The results from Tables 6 and 7 show that Random Forests are a better approach even for the second part of the dataset with fewer examples per class. Considering those results and the probability of an example to fall in the first or second part of the splitted original dataset, final training and validation accuracies were obtained, as may be seen in Table 8.

The best split was in 850, which cut the original dataset in two parts: one with 35464 examples divided into 15 classes and another with 32304 examples divided into 153

Split frequency	Tr. accuracy (%)	Val. accuracy (%)
200	60.81	39.58
300	63.30	41.81
400	65.68	45.77
500	67.87	49.30
550	68.70	49.15
600	69.44	50.16
650	70.03	52.45
750	71.82	53.83
800	72.27	55.04
850	74.57	60.05
900	75.81	61.13
950	76.58	63.43

Table 4. Accuracies for each examples frequency split, with Random Forests being applied to the first part of the dataset with more examples per class

Split frequency	Examples (2nd half)	Classes (2nd half)
200	6001	81
300	10819	109
400	16825	127
500	21254	137
550	22801	140
600	23977	142
650	25859	145
750	27348	147
800	28137	148
850	32304	153
900	34039	155
950	35883	157

Table 5. Number of examples and number of classes in first part of the splitted dataset for each frequency point of split

Split frequency	Tr. accuracy (%)	Val. accuracy (%)
200	41.96	21.39
300	41.59	25.50
400	38.57	23.45
500	36.08	23.39
550	35.97	23.85
600	35.38	24.83
650	35.69	23.39
750	35.54	24.07
800	35.08	24.65
850	35.75	25.65
900	35.15	25.11
950	34.20	23.69

Table 6. Accuracies for each examples frequency split, with SVMs (one vs. one) being applied to the second part of the dataset with fewer examples per class

Split frequency	Tr. accuracy (%)	Val. accuracy (%)
200	72.81	26.48
300	68.56	30.22
400	67.93	28.62
500	65.58	26.99
550	63.75	27.47
600	62.26	27.90
650	62.07	27.82
750	60.80	28.26
800	60.36	28.75
850	60.58	30.15
900	61.08	29.12
950	60.62	28.06

Table 7. Accuracies for each examples frequency split, Random Forests being applied to the second part of the dataset with fewer examples per class

Split frequency	Tr. accuracy (%)	Val. accuracy (%)
200	61.87	38.41
300	64.14	39.96
400	66.24	41.51
500	67.14	42.29
550	67.03	41.85
600	66.90	42.28
650	67.00	43.05
750	67.37	43.51
800	67.32	44.12
850	67.90	45.80
900	68.41	45.05
950	68.13	44.70

Table 8. Final training and validation accuracies for each examples frequency split utilizing Random Forests in each part of the splitted dataset separately

classes. Weighted training and validation accuracies were 67.9% and 45.8%.

5. Discussion

The first approach to the problem was to run a series of classification algorithms for the whole dataset, which resulted in training and validation accuracies of 34.86% and 32.59% for K-Nearest Neighbors, 39.12% and 33.42% for one vs. one SVMs and 58.34% and 36.03% for Random Forests. The best results were given by Random Forests algorithm, showing a significant difference between them, which characterizes as overfitting of the training data even when the algorithm's parameters were tweaked to prevent this issue. Moreover, the validation accuracy was far from a good result.

For the second approach, an ensemble take was applied by dividing the original dataset into two: one with classes

containing a higher number of examples and other with classes containing fewer examples. The first set contained 35464 examples divided into 15 classes and the second 32304 examples divided into 153 classes. Random Forests were utilized in the first one, as it was the algorithm with better performance for the whole dataset, and tends to do better with more examples. For the second, a comparison between performances of one vs. one SVMs and Random Forests was made, with the latter having an overall better performance. Based on the probability of an example to fit in the first or the second set, the final accuracies were calculated for the ensemble: 67.9% and 45.8% for training and validation, respectively, which was a substantial improvement over the first approach but still fell short of even a satisfactory result.

The low accuracies, particularly the validation one, could be explained by a series of factors. The dataset didn't have any truly distinctive feature and almost all of them showed ranges of values that could fit in a great number of classes, reducing the probability of a good classification. Moreover, the imbalance of classes even with the second approach was still present and affected classification, specially for the second portion of the dataset which had a good number of classes with very few examples (less than 50), with the edge case being only 2 examples.

The hypothesis of an overall bad dataset is also enforced when looking to the results of the first portion's classification, where only 15 classes were present, with a great number of examples each (ranging from 862 to 11270), and even in this scenario classification results were modest: 74.57% for training and 60.05% for validation.

6. Conclusion

Even with two different approaches to the problem utilizing powerful multiclass classification algorithms the final results were below expected, which may be explained by the heavily imbalanced dataset and the nature of the features, which doesn't have many distinctions between them to characterize well most classes.